

**** Project - Alternus Vera****

**** Team: code-monkeys****

**** Name: Puja Kawale (012506156)****

**** GitHub:
https://github.com/pujakb/ML/tree/master/News_Classifier****

**** Data preparation and Distillation :****

**** Removing stop words and punchuations****

**** lower case implementation****

**** stemming, lematizing****

**** tokenizing****

**** LDA****

**** LSA****

**** LDA Vs LSA comparision****

**** POS****

**** tSNE****

**** Bokeh clusters****

**** Classifiers on model :****

**** Decision tree, SVM, Naive Bayes, Random Forest, Logistic regression****

Importing the Libraries

```
In [1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
%matplotlib inline
import seaborn as sns
from tqdm import tqdm
from sklearn.feature_extraction.text import CountVectorizer
from sklearn.feature_extraction.text import TfidfTransformer
from sklearn.feature_extraction.text import TfidfVectorizer
from sklearn.pipeline import Pipeline
import nltk
import nltk.corpus
from nltk.tokenize import word_tokenize
import csv
```

Loading the Data

```
In [2]: training = pd.read_csv('C://PUJAMS//machinelearning//csvfiles//liar_dataset//t
rain.tsv',sep='\t')
testing = pd.read_csv('C://PUJAMS//machinelearning//csvfiles//liar_dataset//te
st.tsv',sep='\t')
validation = pd.read_csv('C://PUJAMS//machinelearning//csvfiles//liar_datase
t//valid.tsv',sep='\t')
vocabulary = pd.read_csv (r'C:\PUJAMS\machinelearning\csvfiles\liar_dataset\se
nsation_dictionary.csv')
```

```
In [3]: columns = ['id', 'label', 'statement', 'subjects', 'speaker',
                  'speaker_job', 'state', 'party', 'barely_true_counts',
                  'false_counts', 'half_true_counts', 'mostly_true_counts', 'pants_on_f
                  ire_counts',
                  'context']
training.columns = columns
testing.columns = columns
validation.columns = columns
training.head(5)
```

Out[3]:

	id	label	statement	subjects	speaker	speaker_job	state	
0	10540.json	half-true	When did the decline of coal start? It started...	energy,history,job-accomplishments	scott-surovell	State delegate	Virginia	der
1	324.json	mostly-true	Hillary Clinton agrees with John McCain "by vo...	foreign-policy	barack-obama	President	Illinois	der
2	1123.json	false	Health care reform legislation is likely to ma...	health-care	blog-posting	NaN	NaN	non
3	9028.json	half-true	The economic turnaround started at the end of ...	economy,jobs	charlie-crist	NaN	Florida	der
4	12465.json	true	The Chicago Bears have had more starting quart...	education	robin-vos	Wisconsin Assembly speaker	Wisconsin	repu

```
In [4]: #data integrity check (missing label values)  
#none of the datasets contains missing values therefore no cleaning required  
def data_qualityCheck():  
  
    print("Checking data qualittites...")  
    training.isnull().sum()  
    training.info()  
  
    print("check finished.")  
  
    #below datasets were also used  
    testing.isnull().sum()  
    testing.info()  
  
    validation.isnull().sum()  
    validation.info()  
  
data_qualityCheck()
```

```

Checking data qualittites...
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 10239 entries, 0 to 10238
Data columns (total 14 columns):
id                10239 non-null object
label            10239 non-null object
statement        10239 non-null object
subjects         10237 non-null object
speaker          10237 non-null object
speaker_job      7342 non-null object
state            8031 non-null object
party            10237 non-null object
barely_true_counts 10237 non-null float64
false_counts     10237 non-null float64
half_true_counts 10237 non-null float64
mostly_true_counts 10237 non-null float64
pants_on_fire_counts 10237 non-null float64
context          10137 non-null object
dtypes: float64(5), object(9)
memory usage: 1.1+ MB
check finished.
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1266 entries, 0 to 1265
Data columns (total 14 columns):
id                1266 non-null object
label            1266 non-null object
statement        1266 non-null object
subjects         1266 non-null object
speaker          1266 non-null object
speaker_job      941 non-null object
state            1004 non-null object
party            1266 non-null object
barely_true_counts 1266 non-null int64
false_counts     1266 non-null int64
half_true_counts 1266 non-null int64
mostly_true_counts 1266 non-null int64
pants_on_fire_counts 1266 non-null int64
context          1249 non-null object
dtypes: int64(5), object(9)
memory usage: 138.5+ KB
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1283 entries, 0 to 1282
Data columns (total 14 columns):
id                1283 non-null object
label            1283 non-null object
statement        1283 non-null object
subjects         1283 non-null object
speaker          1283 non-null object
speaker_job      938 non-null object
state            1004 non-null object
party            1283 non-null object
barely_true_counts 1283 non-null int64
false_counts     1283 non-null int64
half_true_counts 1283 non-null int64
mostly_true_counts 1283 non-null int64
pants_on_fire_counts 1283 non-null int64
context          1271 non-null object

```

```
dtypes: int64(5), object(9)  
memory usage: 140.4+ KB
```

Observing the data

```
In [5]: #data observation
def data_obs():
    print("training dataset size:")
    print(training.shape)
    print(training.head(10))

    #below dataset were used for testing and validation purposes
    print(training.shape)
    print(training.head(10))

    print(validation.shape)
    print(validation.head(10))

data_obs()
```

training dataset size:

(10239, 14)

	id	label	statement
0	10540.json	half-true	When did the decline of coal start? It started...
1	324.json	mostly-true	Hillary Clinton agrees with John McCain "by vo...
2	1123.json	false	Health care reform legislation is likely to ma...
3	9028.json	half-true	The economic turnaround started at the end of ...
4	12465.json	true	The Chicago Bears have had more starting quart...
5	2342.json	barely-true	Jim Dunnam has not lived in the district he re...
6	153.json	half-true	I'm the only person on this stage who has work...
7	5602.json	half-true	However, it took \$19.5 million in Oregon Lotte...
8	9741.json	mostly-true	Says GOP primary opponents Glenn Grothman and ...
9	7115.json	mostly-true	For the first time in history, the share of th...

	subjects	speaker
0	energy,history,job-accomplishments	scott-surovell
1	foreign-policy	barack-obama
2	health-care	blog-posting
3	economy,jobs	charlie-crist
4	education	robin-vos
5	candidates-biography	republican-party-texas
6	ethics	barack-obama
7	jobs	oregon-lottery
8	energy,message-machine-2014,voting-record	duey-stroebel
9	elections	robert-menendez

	speaker_job	state	party	barely_true_counts
0	State delegate	Virginia	democrat	0.0
1	President	Illinois	democrat	70.0
2	NaN	NaN	none	7.0
3	NaN	Florida	democrat	15.0
4	Wisconsin Assembly speaker	Wisconsin	republican	0.0
5	NaN	Texas	republican	3.0
6	President	Illinois	democrat	70.0
7	NaN	NaN	organization	0.0
8	State representative	Wisconsin	republican	0.0

9	U.S. Senator	New Jersey	democrat	1.0
	false_counts	half_true_counts	mostly_true_counts	pants_on_fire_counts
\				
0	0.0	1.0	1.0	0.0
1	71.0	160.0	163.0	9.0
2	19.0	3.0	5.0	44.0
3	9.0	20.0	19.0	2.0
4	3.0	2.0	5.0	1.0
5	1.0	1.0	3.0	1.0
6	71.0	160.0	163.0	9.0
7	0.0	1.0	0.0	1.0
8	0.0	0.0	1.0	0.0
9	3.0	1.0	3.0	0.0

context
 0 a floor speech.
 1 Denver
 2 a news release
 3 an interview on CNN
 4 a an online opinion-piece
 5 a press release.
 6 a Democratic debate in Philadelphia, Pa.
 7 a website
 8 an online video
 9 a speech
 (10239, 14)

	id	label	statement
\			
0	10540.json	half-true	When did the decline of coal start? It started...
1	324.json	mostly-true	Hillary Clinton agrees with John McCain "by vo...
2	1123.json	false	Health care reform legislation is likely to ma...
3	9028.json	half-true	The economic turnaround started at the end of ...
4	12465.json	true	The Chicago Bears have had more starting quart...
5	2342.json	barely-true	Jim Dunnam has not lived in the district he re...
6	153.json	half-true	I'm the only person on this stage who has work...
7	5602.json	half-true	However, it took \$19.5 million in Oregon Lotte...
8	9741.json	mostly-true	Says GOP primary opponents Glenn Grothman and ...

9 7115.json mostly-true For the first time in history, the share of th...

	subjects	speaker \
0	energy,history,job-accomplishments	scott-surovell
1	foreign-policy	barack-obama
2	health-care	blog-posting
3	economy,jobs	charlie-crist
4	education	robin-vos
5	candidates-biography	republican-party-texas
6	ethics	barack-obama
7	jobs	oregon-lottery
8	energy,message-machine-2014,voting-record	duey-stroebe
9	elections	robert-menendez

	speaker_job	state	party	barely_true_counts
\				
0	State delegate	Virginia	democrat	0.0
1	President	Illinois	democrat	70.0
2	NaN	NaN	none	7.0
3	NaN	Florida	democrat	15.0
4	Wisconsin Assembly speaker	Wisconsin	republican	0.0
5	NaN	Texas	republican	3.0
6	President	Illinois	democrat	70.0
7	NaN	NaN	organization	0.0
8	State representative	Wisconsin	republican	0.0
9	U.S. Senator	New Jersey	democrat	1.0

	false_counts	half_true_counts	mostly_true_counts	pants_on_fire_counts
\				
0	0.0	1.0	1.0	0.0
1	71.0	160.0	163.0	9.0
2	19.0	3.0	5.0	44.0
3	9.0	20.0	19.0	2.0
4	3.0	2.0	5.0	1.0
5	1.0	1.0	3.0	1.0
6	71.0	160.0	163.0	9.0
7	0.0	1.0	0.0	1.0

8	0.0	0.0	1.0	0.0
9	3.0	1.0	3.0	0.0

context
 0 a floor speech.
 1 Denver
 2 a news release
 3 an interview on CNN
 4 a an online opinion-piece
 5 a press release.
 6 a Democratic debate in Philadelphia, Pa.
 7 a website
 8 an online video
 9 a speech

(1283, 14)

	id	label	statement
\			
0	238.json	pants-fire	When Obama was sworn into office, he DID NOT u...
1	7891.json	false	Says Having organizations parading as being so...
2	8169.json	half-true	Says nearly half of Oregons children are poor.
3	929.json	half-true	On attacks by Republicans that various program...
4	9416.json	false	Says when armed civilians stop mass shootings ...
5	6861.json	true	Says Tennessee is providing millions of dollar...
6	1122.json	false	The health care reform plan would set limits s...
7	13138.json	true	Says Donald Trump started his career back in 1...
8	1880.json	half-true	Bill White has a long history of trying to lim...
9	12803.json	half-true	John McCains chief economic adviser during the...

	subjects	speaker	\
0	obama-birth-certificate,religion	chain-email	
1	campaign-finance,congress,taxes	earl-blumenauer	
2	poverty	jim-francesconi	
3	economy,stimulus	barack-obama	
4	guns	jim-rubens	
5	education,state-budget	andy-berke	
6	health-care	club-growth	
7	candidates-biography,diversity,housing	hillary-clinton	
8	military	republican-party-texas	
9	economy	tim-kaine	

	speaker_job	state	party
\			
0	NaN	NaN	none
1	U.S. representative	Oregon	democrat

2	Member of the State Board of Higher Education	Oregon	none
3	President	Illinois	democrat
4	Small business owner	New Hampshire	republican
5	Lawyer and state senator	Tennessee	democrat
6	NaN	NaN	none
7	Presidential candidate	New York	democrat
8	NaN	Texas	republican
9	U.S. Senator	Virginia	democrat

	barely_true_counts	false_counts	half_true_counts	mostly_true_counts	\
0	11	43	8	5	
1	0	1	1	1	
2	0	1	1	1	
3	70	71	160	163	
4	1	1	0	1	
5	0	0	0	0	
6	4	5	4	2	
7	40	29	69	76	
8	3	1	1	3	
9	8	3	15	15	

	pants_on_fire_counts	context
0	105	NaN
1	0	a U.S. Ways and Means hearing
2	0	an opinion article
3	9	interview with CBS News
4	0	in an interview at gun shop in Hudson, N.H.
5	0	a letter to state Senate education committee c...
6	0	a TV ad
7	7	the first presidential debate
8	1	an e-mail
9	0	a speech at the Democratic National Convention...


```
In [6]: # training=training[['Label', 'statement']]
# training.head(5)
```

In [7]: `training.groupby('label').describe()`

Out[7]:

	barely_true_counts								false_counts		...
	count	mean	std	min	25%	50%	75%	max	count	mean	...
label											
barely-true	1654.0	11.730351	17.605713	1.0	1.0	3.0	13.0	70.0	1654.0	12.422612	...
false	1992.0	11.610944	19.224930	0.0	0.0	2.0	11.0	70.0	1992.0	15.785643	...
half-true	2114.0	11.848628	19.512656	0.0	0.0	2.0	14.0	70.0	2114.0	12.568117	...
mostly-true	1962.0	11.825688	19.999061	0.0	0.0	2.0	14.0	70.0	1962.0	12.280836	...
pants-fire	839.0	11.050060	17.162980	0.0	0.0	5.0	11.0	70.0	839.0	18.162098	...
true	1676.0	10.754773	18.920702	0.0	0.0	1.0	11.0	70.0	1676.0	10.818616	...

6 rows × 40 columns

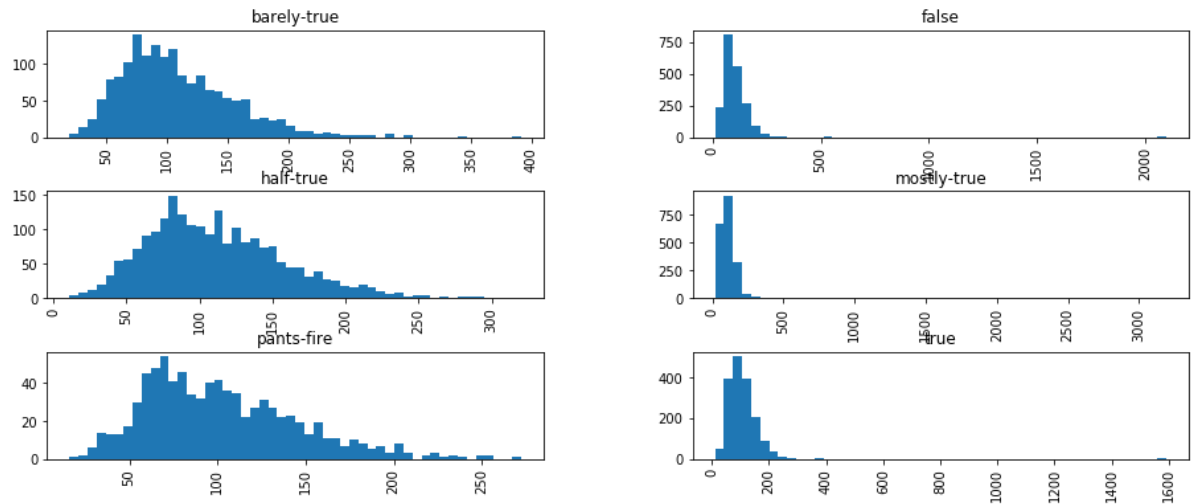


In [8]: `training.iloc[0]["statement"][:len(training.iloc[0]["statement"])]`

Out[8]: 'When did the decline of coal start? It started when natural gas took off tha
t started to begin in (President George W.) Bushs administration.'

```
In [9]: training['length'] = training['statement'].apply(len)
training.hist(column='length',by='label',bins=50, figsize=(15,6))
```

```
Out[9]: array([[<matplotlib.axes._subplots.AxesSubplot object at 0x000000000BE9BE48>,
  <matplotlib.axes._subplots.AxesSubplot object at 0x000000000BEEF668
>],
  [<matplotlib.axes._subplots.AxesSubplot object at 0x000000000BF17978>,
  <matplotlib.axes._subplots.AxesSubplot object at 0x000000000BF41A90
>],
  [<matplotlib.axes._subplots.AxesSubplot object at 0x000000000C07ADA0>,
  <matplotlib.axes._subplots.AxesSubplot object at 0x000000000C07ADD8
>]],
  dtype=object)
```




```
In [10]: #distribution of classes for prediction
def create_distribution(dataFile):

    return sns.countplot(x='label', data=dataFile, palette='husl')


#by calling below we can see that training, test and valid data seems to be fa
ilry evenly distributed between the classes
create_distribution(training)
```

Out[10]: <matplotlib.axes._subplots.AxesSubplot at 0xc23c438>




```
In [11]: create_distribution(testing)
```

Out[11]: <matplotlib.axes._subplots.AxesSubplot at 0xc5729e8>



In [12]: `create_distribution(validation)`

Out[12]: `<matplotlib.axes._subplots.AxesSubplot at 0xc5cc470>`



In [13]: `training.label.unique()`

Out[13]: `array(['half-true', 'mostly-true', 'false', 'true', 'barely-true', 'pants-fire'], dtype=object)`

```
In [15]: # This is in case we want to conduct binary classification
# Our focus is to develop a multi-classification model

# train_news['label']=np.where(train_news['label']=='half-true','true',train_n
ews['label'])
# train_news['label']=np.where(train_news['label']=='mostly-true','true',train
_news['label'])
# train_news['label']=np.where(train_news['label']=='barely-true','false',trai
n_news['label'])
# train_news['label']=np.where(train_news['label']=='pants-fire','false',train
_news['label'])
# train_news['label']=np.where(train_news['label']=='TRUE','true',train_news
['label'])
# train_news['label']=np.where(train_news['label']=='FALSE','false',train_news
['label'])

# plt.title('statement count Vs label')
# train_news.groupby(['label']).size().plot(kind='bar', color='blue')
# plt.xlabel('count')
# plt.show()
```



```
In [16]: import re
import numpy as np
import matplotlib.pyplot as plt
import pandas as pd
import nltk
nltk.download('stopwords')
from nltk.corpus import stopwords
from nltk.stem.porter import PorterStemmer
from nltk.tokenize import word_tokenize
nltk.download('punkt')

def TextCleansing(txt):
    rem = re.sub('[^a-zA-Z]', ' ', txt)
    rem = rem.lower()
    rem = rem.split()
    ps = PorterStemmer()
    rem = [ps.stem(word) for word in rem if not word in set(stopwords.words('english'))]
    rem = ' '.join(rem)
    return rem
```

```
[nltk_data] Downloading package stopwords to
[nltk_data]   C:\Users\PB\AppData\Roaming\nltk_data...
[nltk_data]   Package stopwords is already up-to-date!
[nltk_data] Downloading package punkt to
[nltk_data]   C:\Users\PB\AppData\Roaming\nltk_data...
[nltk_data]   Package punkt is already up-to-date!
```

```
In [17]: stop_words=set(stopwords.words('english'))
print(stop_words)
```

```
{'to', 'of', 'once', 'he', 'into', 't', 'ours', 'doing', "weren't", "should've", 'if', 'through', 'yours', 'am', 'during', 'out', "aren't", "wouldn't", 'd', 'id', "won't", 'wouldn', 'all', 'they', 'below', 'do', 'few', 'been', 'can', 'isn', 'by', 'against', 'i', 'why', 've', 'whom', 'now', 'himself', 'and', 'o', 'r', 'we', 'in', 'couldn', "mightn't", 'you', 'being', 'own', 'other', 'aren', 'who', 'was', 'are', 'shouldn', 'mightn', "don't", 'what', 'for', 'mustn', 'd', 'own', "you'll", 'a', 'their', 'than', 'most', "doesn't", 'its', 'very', 'wer', 'e', 'won', "couldn't", 'some', 'about', 'same', 'on', 'your', 'herself', 'bu', 't', 'ma', 'while', 'when', 'have', "shouldn't", "hadn't", 'more', 'my', "yo", 'u'd', 'over', 'haven', 'not', "hasn't", "you're", 'be', 'is', 'just', 'her', "she's", 'from', 'him', 'because', 'above', 'there', 'those', 'which', 'sha', 'n', "wasn't", 'it', 'll', 'his', "haven't", 'further', 'hers', 'weren', 'unti', 'l', 'ain', 'our', 'doesn', 'before', 'd', "mustn't", 'only', 'both', 'me', "d", 'idn't', 'after', 'this', 'between', 'hasn', 'myself', "it's", 'has', 'y', 'to', 'o', "shan't", 'themselves', 'with', 'again', 'itself', 'under', "that'll", 's', 'he', 'didn', 'the', 'an', 'then', 'don', 'that', 'up', 're', 'such', 'where', 'hadn', 'as', "isn't", 'ourselves', 'will', 'o', 'wasn', 'how', 'each', 'an', 'y', "needn't", 'had', 'needn', 'at', 'yourself', "you've", 'yourselves', 'n', 'o', 'nor', 'theirs', 'm', 'here', 'does', 'these', 'them', 'should', 'havin', 'g', 'so', 'off', 's'}
```

```
In [19]: training_head = training[['statement']]
training_head['index'] = training_head.index
documents = training_head
```

C:\Users\PB\Anaconda3\lib\site-packages\ipykernel_launcher.py:2: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: <http://pandas.pydata.org/pandas-docs/stable/indexing.html#indexing-view-versus-copy>

```
In [20]: print(len(documents))
print(documents[:5])
```

10239

	statement	index
0	When did the decline of coal start? It started...	0
1	Hillary Clinton agrees with John McCain "by vo...	1
2	Health care reform legislation is likely to ma...	2
3	The economic turnaround started at the end of ...	3
4	The Chicago Bears have had more starting quart...	4

Data Preprocessing

```
In [21]: import gensim
from gensim.utils import simple_preprocess
from gensim.parsing.preprocessing import STOPWORDS
from nltk.stem import WordNetLemmatizer, SnowballStemmer
from nltk.stem.porter import *
import numpy as np
np.random.seed(2018)
import nltk
nltk.download('wordnet')
```

C:\Users\PB\Anaconda3\lib\site-packages\gensim\utils.py:1212: UserWarning: detected Windows; aliasing chunkize to chunkize_serial
warnings.warn("detected Windows; aliasing chunkize to chunkize_serial")

[nltk_data] Downloading package wordnet to
[nltk_data] C:\Users\PB\AppData\Roaming\nltk_data...
[nltk_data] Package wordnet is already up-to-date!

Out[21]: True

```
In [22]: from nltk.stem import PorterStemmer, WordNetLemmatizer
def text_preprocessing(df_base):
    df=df_base.copy()
    # lowercase the text
    df['statement']=df['statement'].str.lower()
    # word tokenization
    df['statement']=df['statement'].map(lambda x: nltk.word_tokenize(x))
    # remove stop words and non alphanumeric charaters
    df['statement']=df['statement'].map(lambda x: [w for w in x if (not w in s
top_words) and w.isalpha()]])
    # lemmatization
    wordnet_lemmatizer = WordNetLemmatizer()
    df['statement']=df['statement'].map(lambda x: [ wordnet_lemmatizer.lemmati
ze(w) for w in x])
    # stemming
    porter = PorterStemmer()
    df['statement']=df['statement'].map(lambda x: [porter.stem(w) for w in x]
)
    return df
```

```
In [23]: processed =text_preprocessing(training)
processed.head()
```

Out[23]:

	id	label	statement	subjects	speaker	speaker_job	state	
0	10540.json	half-true	[declin, coal, start, start, natur, ga, took, ...	energy,history,job-accomplishments	scott-surovell	State delegate	Virginia	de
1	324.json	mostly-true	[hillari, clinton, agre, john, mccain, vote, g...	foreign-policy	barack-obama	President	Illinois	de
2	1123.json	false	[health, care, reform, legisl, like, mandat, f...	health-care	blog-posting	NaN	NaN	nc
3	9028.json	half-true	[econom, turnaround, start, end, term]	economy,jobs	charlie-crist	NaN	Florida	de
4	12465.json	true	[chicago, bear, start, quarterback, last, year...	education	robin-vos	Wisconsin Assembly speaker	Wisconsin	re

Word2Vec

Word2vec is a two-layer neural net that processes text. Its input is a text corpus and its output is a set of vectors: feature vectors for words in that corpus. It turns text into a numerical form that deep nets can understand

```
In [24]: import gensim
from gensim.models.word2vec import Word2Vec
from gensim.models.doc2vec import TaggedDocument
from sklearn.model_selection import train_test_split
import warnings
warnings.filterwarnings('ignore')
```

```
In [25]: LabeledSentence = gensim.models.doc2vec.LabeledSentence
```

```
In [26]: def labelled_sentences(articles, label_type):
    labelledSentences = []
    for i,_d in enumerate(articles):
        labelledSentences.append(LabeledSentence(_d, label_type[i]))
    return labelledSentences
```

```
In [27]: processed_labelled=labelled_sentences(processed['statement'],processed['label'
])
processed_labelled[1]
```

```
Out[27]: LabeledSentence(words=['hillari', 'clinton', 'agre', 'john', 'mccain', 'vot
e', 'give', 'georg', 'bush', 'benefit', 'doubt', 'iran'], tags='mostly-true')
```

```
In [28]: tagged = [TaggedDocument(words=_d, tags=[str(i)]) for i, _d in enumerate(proce
ssed['statement'])]
```

```
In [29]: tagged_words=[x.words for x in tagged]
tagged_words_1D=[]
for row in range(len(tagged_words)):
    for col in range(len(tagged_words[row])):
        tagged_words.append(tagged_words[row][col])

n_dim=300
w2v = Word2Vec(size=n_dim, min_count=0)
w2v.build_vocab(tagged_words)
w2v.train(tagged_words,total_examples=w2v.corpus_count,epochs=w2v.epochs)
```

```
Out[29]: (1043861, 3317680)
```

In [30]: `w2v.most_similar('war')`

Out[30]: `[('rapid', 0.9978175759315491),
('longest', 0.9975180625915527),
('fundrais', 0.9974884986877441),
('incid', 0.997333824634552),
('extrem', 0.9971774816513062),
('watch', 0.9971362948417664),
('jump', 0.9971160292625427),
('gasolin', 0.9970066547393799),
('happen', 0.996982216835022),
('document', 0.9969084858894348)]`

In [31]: `processed_statements=processed[['statement','label']]
processed_statements_vectorized=processed_statements['statement'].map(lambda x
: [w2v[w] for w in x])

processed_statements['statement']=processed_statements_vectorized

from sklearn import preprocessing
le = preprocessing.LabelEncoder()
List(le.inverse_transform(processed_statements['label']))
processed_statements['label']=le.fit_transform(processed_statements['label'])

processed_statements.head()`

Out[31]:

	statement	label
0	[[0.08882012, 0.052362617, 0.044050723, 0.0263...	2
1	[[0.23452127, -0.103835195, -0.14599872, 0.228...	3
2	[[0.73109084, 0.56986696, -0.44808704, -0.1760...	1
3	[[0.22355153, 0.084662944, 0.04774518, 0.07887...	2
4	[[0.09246858, 0.02514211, 0.0040849606, 0.0322...	5

```
In [32]: def document_vector(doc):  
         return np.mean(doc, axis=0)  
  
         for row_idx in range(len(processed_statements['statement'])):  
             for col_idx in range(len(processed_statements['statement'][row_idx])):  
                 processed_statements['statement'][row_idx][col_idx]=document_vector(pr  
ocessed_statements['statement'][row_idx][col_idx])  
  
         #df_train_statements['statement']=df_train_statements['statement'].map(documen  
t_vector)  
         processed_statements.drop(index=4497,inplace=True)  
         processed_statements.reset_index(drop=True)
```

Out[32]:

	statement	label
0	[0.0015531193, 0.0030009032, 0.0048259343, 0.0...	2
1	[0.005818305, 0.010304897, 0.0036140848, 0.007...	3
2	[0.0039481856, 0.018210806, 0.00598304, 0.0067...	1
3	[0.0058103004, 0.001316087, 0.0048259343, 0.00...	2
4	[0.0027394965, 0.002334294, 0.0048259343, -2.0...	5
5	[0.0051740645, 8.1900405e-05, 0.007307156, 0.0...	0
6	[0.0063952627, 0.003622454, 0.0065158713, 0.00...	2
7	[0.0009564376, 0.0042650383, 0.006832938, 0.00...	2
8	[0.0073591056, 0.004535676, 0.0032518224, 0.00...	3
9	[0.004452211, 0.0056764404, 0.0058674715, 0.00...	3
10	[0.00523225, 0.006501026, 0.006832938, 0.00713...	2
11	[0.0046875635, 0.0063856672, 0.006251499, 0.00...	1
12	[0.005886208, 0.00011750842, 0.0055083116, 0.0...	3
13	[0.0020381382, 0.018210806, 0.0046027345, 0.00...	0
14	[0.0048613134, 0.0055349213, 0.0010576568, 0.0...	2
15	[0.002692792, 0.0048573655, 0.0063228006, 0.00...	5
16	[0.0025115563, 0.0054791453, 0.0026679575, 0.0...	0
17	[0.007653927, 0.009440354, 8.842249e-06, 0.002...	2
18	[0.005894712, 0.007918409, 0.005161361, 0.0018...	3
19	[0.0074008037, 0.0042240904, 0.007807286, 0.00...	1
20	[0.019988174, 0.013930908, 0.0056922636, 0.004...	3
21	[0.0055046654, 0.007970379, 0.003117077, 0.006...	3
22	[0.0073591056, 0.0065401616, 0.006110296, 0.00...	2
23	[0.0073591056, 0.0046875635, 0.0063856672, 0.0...	0
24	[0.0060428553, 0.00692223, 0.00022820162, 0.00...	1
25	[0.0025243398, 0.006544179, 0.007132008, 0.006...	3
26	[0.00528491, 0.0038272997, 0.005425787, 0.0061...	2
27	[0.007084581, 0.0053566354, 0.0017373619, 0.00...	1
28	[0.0011862061, 0.0064320765, 0.0042735455, 0.0...	3
29	[0.0073591056, 0.0037955025, 0.0019614815, 0.0...	5
...
10208	[0.00523225, 0.0020381382, 0.018210806, 0.0046...	5

	statement	label
10209	[0.0043656705, 0.005894712, 0.0040787444, 0.00...	5
10210	[0.0053095864, 0.0073591056, 0.00654957, 0.006...	2
10211	[0.0049790507, 0.0012464045, 0.0074034277, 0.0...	0
10212	[0.0048613134, 0.007970379, 0.0039481856, 0.01...	1
10213	[0.018964062, -0.0027135627, 0.004326349, 0.00...	0
10214	[0.0023030671, 0.0054027964, 0.0053100125, 0.0...	1
10215	[0.003786529, 0.018964062, -0.0027135627, 0.00...	2
10216	[0.0043705986, 0.008069582, 0.004581774, 0.003...	1
10217	[0.007304689, 0.007270716, 0.006722042, 0.0034...	2
10218	[0.0073591056, 0.0046875635, 0.0063856672, 0.0...	3
10219	[0.0062789354, 0.0053005423, 0.0064634476, 0.0...	1
10220	[0.004452211, 0.0056764404, 0.00523225, 0.0018...	5
10221	[0.0073591056, 0.00528491, 0.0038272997, 0.005...	3
10222	[0.013930908, 0.0064594317, 0.0042387014, 0.00...	0
10223	[0.00654957, 0.0071791466, 0.006127811, 0.0068...	2
10224	[0.00024483484, 0.001080337, 0.0016555116, 0.0...	3
10225	[0.0040249764, 0.0032681953, 0.005939055, 0.00...	4
10226	[0.00600365, 0.00079215714, 0.0013603992, 0.00...	0
10227	[0.0038574457, 0.00045269995, 0.0059986347, 0....	3
10228	[0.005329299, 0.008814389, 0.0025323606, 0.003...	0
10229	[0.007026281, 0.013930908, 0.0030334643, 0.009...	2
10230	[0.0010748486, 0.0067132316, 0.0023030671, 0.0...	2
10231	[0.007120289, 0.0010893724, 0.0052630813, 0.00...	4
10232	[0.0073080175, 0.005329299, 0.008814389, 0.006...	2
10233	[0.00342068, 0.0074034277, 0.00023298124, 0.00...	3
10234	[0.008069582, 0.0048294296, 0.0071791466, 0.00...	3
10235	[0.0073591056, 0.0006689055, 0.004880984, 0.00...	2
10236	[0.0011983047, 0.0010808384, 0.0012000249, 0.0...	1
10237	[0.008196405, 0.004830276, 0.0021933161, 0.000...	4

10238 rows × 2 columns


```

In [33]: texts=[]
         for x in processed_statements['statement'].values:
             texts.append(x)
         X=pd.DataFrame(texts)
         X=X.fillna(0)
         y=processed_statements[['label']]
         X_train,X_test,y_train,y_test=train_test_split(X, y,test_size = .2, random_state = 1)
         X_train.head()

```

Out[33]:

	0	1	2	3	4	5	6	7	
8892	0.005431	0.001999	0.007128	0.002360	0.006550	0.007985	0.004348	0.007403	0.
2933	0.007359	0.009108	0.013931	0.003614	0.002833	0.002692	0.013931	0.001240	0.
3933	0.003033	0.007970	0.005704	0.006502	0.004355	0.006663	0.002504	0.001453	0.
3406	0.007359	-0.003100	0.018964	0.001833	0.003709	0.002651	0.000533	0.005263	0.
7210	0.006073	0.005896	0.003588	0.001695	0.006648	0.005372	0.001766	0.000000	0.

5 rows × 245 columns



Classifiers :

- 1) Logistic Regression
- 2) Decision Tree
- 3) GNB - Gaussian Naive Bayes
- 4) Random Forest Classifier

```
In [34]: from sklearn import metrics
from sklearn.metrics import accuracy_score

from sklearn.model_selection import train_test_split
from sklearn import linear_model

lm = linear_model.LogisticRegression(verbose=1)
model = lm.fit(X_train, y_train)
print (model)
predictions = lm.predict(X_test)

print ("Score:", model.score(X_test, y_test))

print(metrics.classification_report(y_test,predictions))

[LibLinear]LogisticRegression(C=1.0, class_weight=None, dual=False, fit_intercept=True,
    intercept_scaling=1, max_iter=100, multi_class='ovr', n_jobs=1,
    penalty='l2', random_state=None, solver='liblinear', tol=0.0001,
    verbose=1, warm_start=False)
Score: 0.2109375
```

	precision	recall	f1-score	support
0	0.00	0.00	0.00	309
1	0.00	0.00	0.00	396
2	0.21	1.00	0.35	432
3	0.00	0.00	0.00	392
4	0.00	0.00	0.00	171
5	0.00	0.00	0.00	348
avg / total	0.04	0.21	0.07	2048

```
In [35]: from sklearn.model_selection import train_test_split
from sklearn.tree import DecisionTreeClassifier
from sklearn.linear_model import LogisticRegression

clf_gini = DecisionTreeClassifier(criterion = "gini", random_state = 100,
                                max_depth=3, min_samples_leaf=5)
clf_gini.fit(X_train, y_train)
dec_pred = clf_gini.predict(X_test)
print ("Score:", clf_gini.score(X_test, y_test))

print(metrics.classification_report(y_test,dec_pred))
```

Score: 0.2099609375

	precision	recall	f1-score	support
0	0.00	0.00	0.00	309
1	0.26	0.12	0.16	396
2	0.23	0.23	0.23	432
3	0.20	0.73	0.31	392
4	0.00	0.00	0.00	171
5	0.00	0.00	0.00	348
avg / total	0.14	0.21	0.14	2048

```
In [36]: import numpy as np
from sklearn import metrics
from sklearn.metrics import accuracy_score

from sklearn.naive_bayes import GaussianNB
gnb = GaussianNB().fit(X_train, y_train)
gnb_pred = gnb.predict(X_test)

print ("Score:", gnb.score(X_test, y_test))
print(metrics.classification_report(y_test,gnb_pred))
```

Score: 0.08935546875

	precision	recall	f1-score	support
0	0.25	0.00	0.01	309
1	0.00	0.00	0.00	396
2	0.32	0.03	0.05	432
3	0.00	0.00	0.00	392
4	0.09	1.00	0.16	171
5	0.00	0.00	0.00	348
avg / total	0.11	0.09	0.02	2048

```
In [37]: from sklearn.ensemble import RandomForestClassifier
num_trees=200

rf=RandomForestClassifier(n_estimators=num_trees)
rf.fit(X_train, y_train)
rf_pred=rf.predict(X_test)

print ("Score:", rf.score(X_test, y_test))
print(metrics.classification_report(y_test,rf_pred))
```

Score: 0.2177734375

	precision	recall	f1-score	support
0	0.19	0.16	0.17	309
1	0.23	0.25	0.24	396
2	0.22	0.35	0.27	432
3	0.22	0.26	0.24	392
4	0.14	0.01	0.02	171
5	0.21	0.12	0.15	348
avg / total	0.21	0.22	0.20	2048

From above we can conclude that Logistic regression and Decision Tree classifier are performing better than the other classifiers.

Visualization:

A popular method for visualizing document similarity is to use t-distributed stochastic neighbor embedding, t-SNE.

1) Load the corpus and vectorize the text using TF-IDF.

2) Once the corpus is vectorized we visualize it, showing the distribution of classes.

reference: <http://www.scikit-yb.org/en/latest/api/text/tsne.html> (<http://www.scikit-yb.org/en/latest/api/text/tsne.html>)

```
In [38]: # Import Libraries
from sklearn.manifold import TSNE
import matplotlib as mpl
import matplotlib.pyplot as plt
warnings.filterwarnings("ignore",category=RuntimeWarning)

# List of vocabulary.

vocab = list(w2v.wv.vocab)

# index vector values by corresponding vocab list

X = w2v[vocab]

print("Total Number of Vocab:", len(X))
print()
print(X[0][:10])

# Visualize only 100 words.

tsne = TSNE(n_components = 2)
X_tsne = tsne.fit_transform(X[:100,:])

df = pd.DataFrame(X_tsne, index = vocab[:100], columns = ['X','Y'])
df.head()


fig = plt.figure()
fig.set_size_inches(30,20)

ax = fig.add_subplot(1,1,1)
ax.scatter(df['X'], df['Y'])

# Put the Label on each point.
for word, pos in df.iterrows():
    ax.annotate(word, pos, fontsize = 20)
plt.show()
```

Total Number of Vocab: 7539

[0.08882012 0.05236262 0.04405072 0.02633099 0.08770611 -0.09703708
0.02267416 -0.05655846 -0.10878675 0.05518287]



```
In [44]: import numpy as np
import pandas as pd
from IPython.display import display
from tqdm import tqdm
from collections import Counter
import ast

import matplotlib.pyplot as plt
import matplotlib.mlab as mlab
import seaborn as sb

%matplotlib inline

raw_data = training

reindexed_data = raw_data['statement']
reindexed_data.index = raw_data['label']

display(raw_data.head())
```

	id	label	statement	subjects	speaker	speaker_job	state	
0	10540.json	half-true	When did the decline of coal start? It started...	energy,history,job-accomplishments	scott-surovell	State delegate	Virginia	der
1	324.json	mostly-true	Hillary Clinton agrees with John McCain "by vo...	foreign-policy	barack-obama	President	Illinois	der
2	1123.json	false	Health care reform legislation is likely to ma...	health-care	blog-posting	NaN	NaN	non
3	9028.json	half-true	The economic turnaround started at the end of ...	economy,jobs	charlie-crist	NaN	Florida	der
4	12465.json	true	The Chicago Bears have had more starting quart...	education	robin-vos	Wisconsin Assembly speaker	Wisconsin	rept

Developing list of the top words used in the statements of Liar Liar dataset, . Stop words are removed from the dataset to avoid any trivial conjunctions, prepositions, etc.


```
In [45]: # Define helper functions
def get_top_n_words(n_top_words, count_vectorizer, text_data):
    '''returns a tuple of the top n words in a sample and their accompanying counts, given a CountVectorizer object and text sample'''
    vectorized_headlines = count_vectorizer.fit_transform(text_data.as_matrix())

    vectorized_total = np.sum(vectorized_headlines, axis=0)
    word_indices = np.flip(np.argsort(vectorized_total)[0,:], 1)
    word_values = np.flip(np.sort(vectorized_total)[0,:], 1)

    word_vectors = np.zeros((n_top_words, vectorized_headlines.shape[1]))
    for i in range(n_top_words):
        word_vectors[i, word_indices[0,i]] = 1

    words = [word[0].encode('ascii').decode('utf-8') for word in count_vectorizer.inverse_transform(word_vectors)]

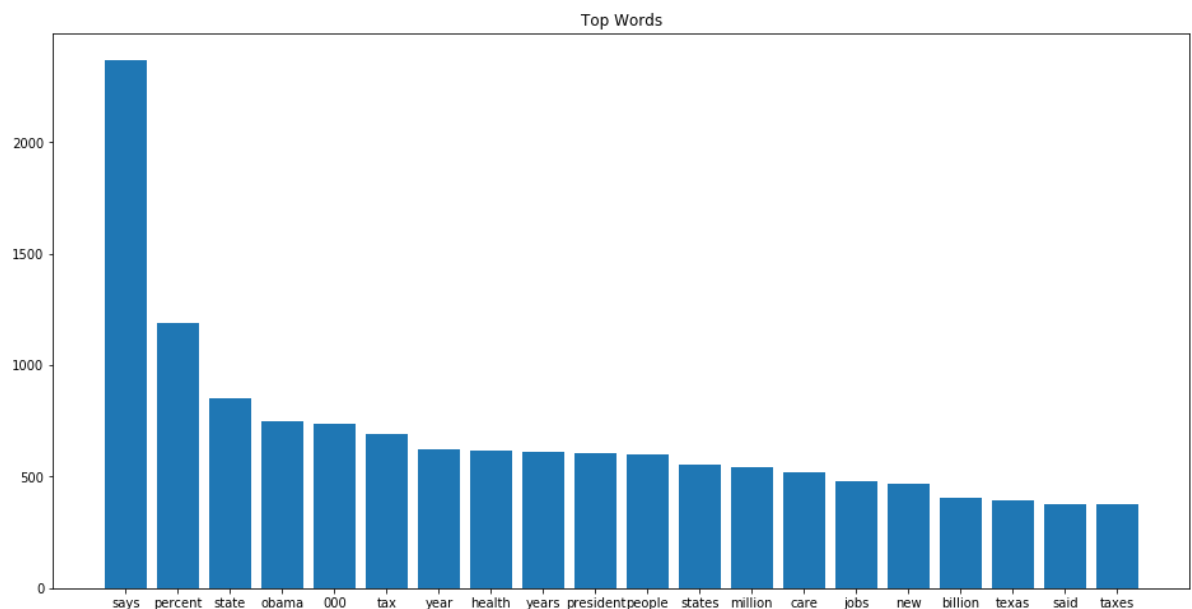
    return (words, word_values[0,:n_top_words].tolist()[0])
```

```
In [46]: from sklearn.feature_extraction.text import CountVectorizer

count_vectorizer = CountVectorizer(stop_words='english')
words, word_values = get_top_n_words(n_top_words=20, count_vectorizer=count_vectorizer, text_data=reindexed_data)

fig, ax = plt.subplots(figsize=(16,8))
ax.bar(range(len(words)), word_values)
ax.set_xticks(range(len(words)))
ax.set_xticklabels(words)
ax.set_title('Top Words')
```

Out[46]: Text(0.5,1,'Top Words')



Part Of Speech Tagging

Below is the histogram of statement word lengths, and (POS) part-of-speech tagging to understand the types of words used across the corpus. This requires first converting all headline strings to TextBlobs and calling the `pos_tags` method on each, yielding a list of tagged words for each headline. A complete list of such word tags is available [here](#).

```
In [47]: from textblob import TextBlob

while True:
    try:
        tagged_headlines = pd.read_csv('abcnews-pos-tagged.csv', index_col=0)
        word_counts = []
        pos_counts = {}

        for headline in tagged_headlines[u'tags']:
            headline = ast.literal_eval(headline)
            word_counts.append(len(headline))
            for tag in headline:
                if tag[1] in pos_counts:
                    pos_counts[tag[1]] += 1
                else:
                    pos_counts[tag[1]] = 1

    except IOError:
        tagged_headlines = [TextBlob(reindexed_data[i]).pos_tags for i in range(reindexed_data.shape[0])]

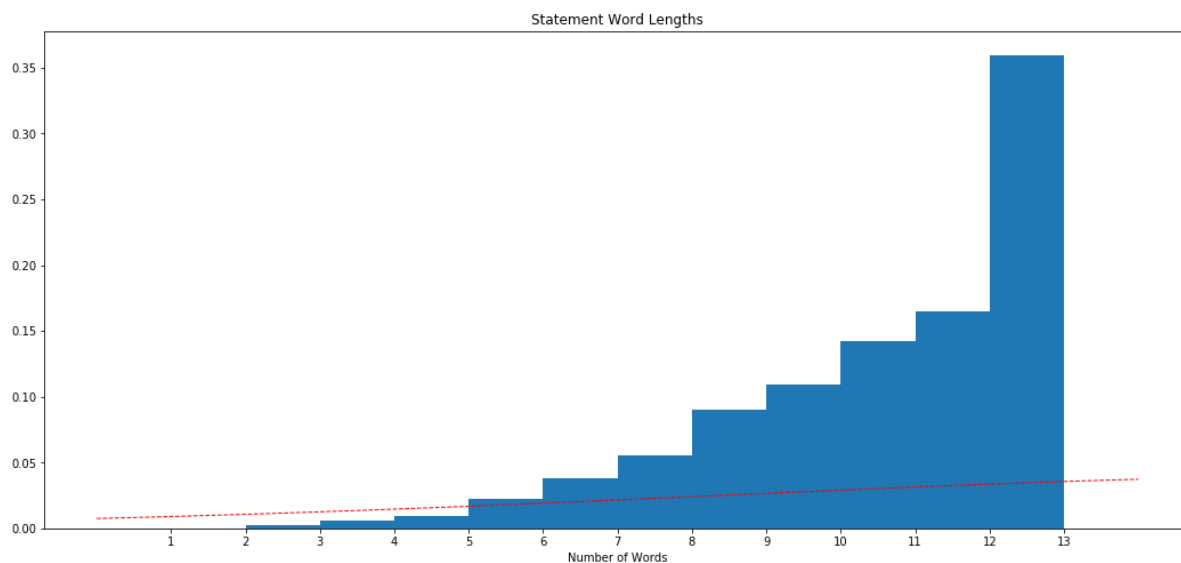
        tagged_headlines = pd.DataFrame({'tags':tagged_headlines})
        tagged_headlines.to_csv('abcnews-pos-tagged.csv')
        continue
    break

print('Total number of words: ', np.sum(word_counts))
print('Mean number of words per headline: ', np.mean(word_counts))
```

Total number of words: 184817

Mean number of words per headline: 18.05029788065241

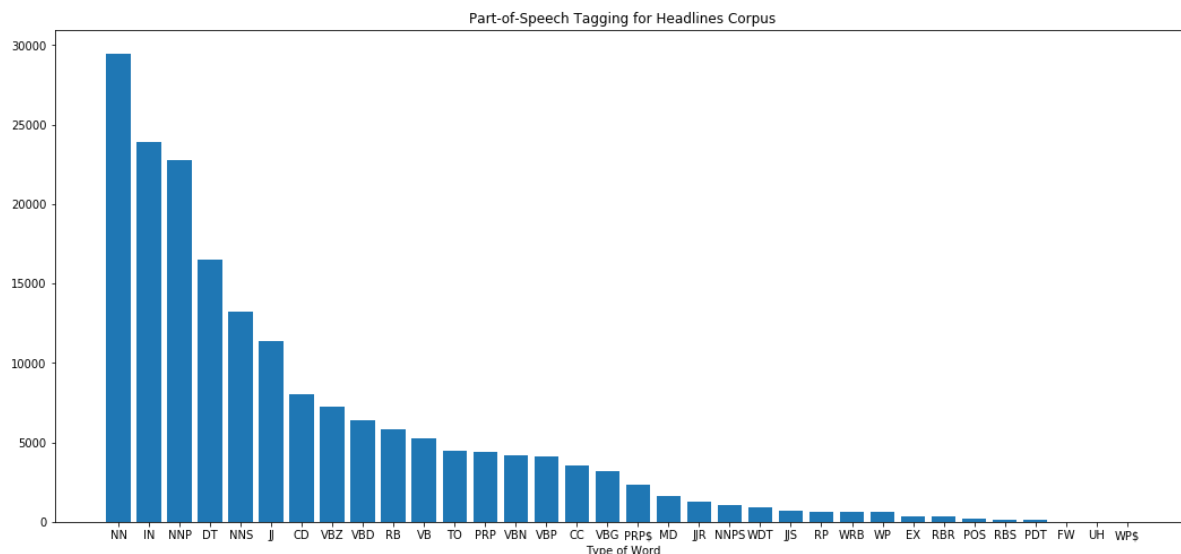
```
In [48]: fig, ax = plt.subplots(figsize=(18,8))
ax.hist(word_counts, bins=range(1,14), normed=1)
ax.set_title('Statement Word Lengths')
ax.set_xticks(range(1,14))
ax.set_xlabel('Number of Words')
y = mlab.normpdf( np.linspace(0,14,50), np.mean(word_counts), np.std(word_counts))
l = ax.plot(np.linspace(0,14,50), y, 'r--', linewidth=1)
```



```
In [49]: pos_sorted_types = sorted(pos_counts, key=pos_counts.__getitem__, reverse=True)
pos_sorted_counts = sorted(pos_counts.values(), reverse=True)

fig, ax = plt.subplots(figsize=(18,8))
ax.bar(range(len(pos_counts)), pos_sorted_counts)
ax.set_xticks(range(len(pos_counts)))
ax.set_xticklabels(pos_sorted_types)
ax.set_title('Part-of-Speech Tagging for Headlines Corpus')
ax.set_xlabel('Type of Word')
```

Out[49]: Text(0.5,0,'Type of Word')



Topic Modelling

Now, let's apply clustering algorithm to the statements in order to study the focus of topic of all the news contained. I had experimented this with a small subsample of the dataset in order to determine which of the two potential clustering algorithms is most appropriate.

Preprocessing - Feature Construction; where I would take the sample of text headlines and represent them in feature space. It's nothing but, converting each string to a numerical vector. To do this I used, CountVectorizer from SKLearn, which yields an $n \times K$ document-term matrix where K is the number of distinct words across the n headlines in our sample with less stop words and with a limit of `max_features`.

```
In [50]: small_count_vectorizer = CountVectorizer(stop_words='english', max_features=40000)
small_text_sample = reindexed_data.sample(n=10000, random_state=0).as_matrix()

print('Headline before vectorization: ', small_text_sample[123])

small_document_term_matrix = small_count_vectorizer.fit_transform(small_text_sample)

print('Headline after vectorization: \n', small_document_term_matrix[123])
```

Headline before vectorization: (U.S. Reps.) Paul Ryan, Sean Duffy and Reid Ribble are shutting down town hall meetings, or making their constituents pay to attend them.

Headline after vectorization:

(0, 1232)	1
(0, 2644)	1
(0, 6811)	1
(0, 5028)	1
(0, 10809)	1
(0, 9729)	1
(0, 9117)	1
(0, 8856)	1
(0, 3668)	1
(0, 9488)	1
(0, 9288)	1
(0, 8979)	1
(0, 7812)	1
(0, 7823)	1
(0, 6599)	1

Now I would apply the clustering algorithm.

1) Latent Semantic Analysis, OR

2) Latent Dirichlet Allocation.

Both will take document-term matrix as input and yield an $n \times N$ topic matrix as output, where N is the number of topic categories, a parameter we can say, which would be 8 in this case.

```
In [51]: n_topics = 8
```

Latent Semantic Analysis

LSA is a technique in natural language processing, in particular distributional semantics, of analyzing relationships between a set of documents and the terms they contain by producing a set of concepts related to the documents and terms. LSA assumes that words that are close in meaning will occur in similar pieces of text

```
In [52]: # Taking the argmax of each headline in this topic matrix will give the predicted topics of each headline in the sample.
# We can then sort these into counts of each topic.
# Define helper functions
def get_keys(topic_matrix):
    '''returns an integer list of predicted topic categories for a given topic matrix'''
    keys = []
    for i in range(topic_matrix.shape[0]):
        keys.append(topic_matrix[i].argmax())
    return keys

def keys_to_counts(keys):
    '''returns a tuple of topic categories and their accompanying magnitudes for a given list of keys'''
    count_pairs = Counter(keys).items()
    categories = [pair[0] for pair in count_pairs]
    counts = [pair[1] for pair in count_pairs]
    return (categories, counts)
```

```
In [53]: from sklearn.decomposition import TruncatedSVD

lsa_model = TruncatedSVD(n_components=n_topics)
lsa_topic_matrix = lsa_model.fit_transform(small_document_term_matrix)

lsa_keys = get_keys(lsa_topic_matrix)
lsa_categories, lsa_counts = keys_to_counts(lsa_keys)
```

```
In [54]: # Define helper functions
def get_top_n_words(n, keys, document_term_matrix, count_vectorizer):
    '''returns a list of n_topic strings, where each string contains the n most common words in a predicted category, in order'''
    top_word_indices = []
    for topic in range(n_topics):
        temp_vector_sum = 0
        for i in range(len(keys)):
            if keys[i] == topic:
                temp_vector_sum += document_term_matrix[i]
        temp_vector_sum = temp_vector_sum.toarray()
        top_n_word_indices = np.flip(np.argsort(temp_vector_sum)[0][-n:],0)
        top_word_indices.append(top_n_word_indices)
    top_words = []
    for topic in top_word_indices:
        topic_words = []
        for index in topic:
            temp_word_vector = np.zeros((1,document_term_matrix.shape[1]))
            temp_word_vector[:,index] = 1
            the_word = count_vectorizer.inverse_transform(temp_word_vector)[0][0]
            topic_words.append(the_word.encode('ascii').decode('utf-8'))
        top_words.append(" ".join(topic_words))
    return top_words
```

```
In [55]: top_n_words_lsa = get_top_n_words(10, lsa_keys, small_document_term_matrix, small_count_vectorizer)

for i in range(len(top_n_words_lsa)):
    print("Topic {}: ".format(i), top_n_words_lsa[i])
```


```
Topic 0: says years people texas voted said new million state obama
Topic 1: percent rate year people 40 10 years unemployment state income
Topic 2: health care 000 million insurance law year people plan government
Topic 3: state 000 year jobs new georgia 10 million billion dollars
Topic 4: obama president 000 jobs barack year years administration bush created
Topic 5: state budget wisconsin florida republican years billion rhode texas georgia
Topic 6: tax taxes plan income cuts billion increase cut middle class
Topic 7: states united country world republican million people countries half senate
```

Thus now I converted the small subset of statements into a list of predicted topic categories, where each category is characterised by its most frequent words.

```
In [56]: top_3_words = get_top_n_words(3, lsa_keys, small_document_term_matrix, small_count_vectorizer)
labels = ['Topic {}: \n'.format(i) + top_3_words[i] for i in lsa_categories]

fig, ax = plt.subplots(figsize=(16,8))
ax.bar(lsa_categories, lsa_counts)
ax.set_xticks(lsa_categories)
ax.set_xticklabels(labels)
ax.set_title('LSA Topic Category Counts')
```

Out[56]: Text(0.5,1,'LSA Topic Category Counts')



In order to properly differentiate LSA with LDA , lets try using dimensionality-reduction technique called t -SNE, which will also serve to better illuminate the success of the clustering process.

```
In [57]: from sklearn.manifold import TSNE

tsne_lsa_model = TSNE(n_components=2, perplexity=50, learning_rate=100,
                      n_iter=2000, verbose=1, random_state=0, angle=0.75)
tsne_lsa_vectors = tsne_lsa_model.fit_transform(lsa_topic_matrix)

[t-SNE] Computing 151 nearest neighbors...
[t-SNE] Indexed 10000 samples in 0.023s...
[t-SNE] Computed neighbors for 10000 samples in 1.360s...
[t-SNE] Computed conditional probabilities for sample 1000 / 10000
[t-SNE] Computed conditional probabilities for sample 2000 / 10000
[t-SNE] Computed conditional probabilities for sample 3000 / 10000
[t-SNE] Computed conditional probabilities for sample 4000 / 10000
[t-SNE] Computed conditional probabilities for sample 5000 / 10000
[t-SNE] Computed conditional probabilities for sample 6000 / 10000
[t-SNE] Computed conditional probabilities for sample 7000 / 10000
[t-SNE] Computed conditional probabilities for sample 8000 / 10000
[t-SNE] Computed conditional probabilities for sample 9000 / 10000
[t-SNE] Computed conditional probabilities for sample 10000 / 10000
[t-SNE] Mean sigma: 0.021052
[t-SNE] KL divergence after 250 iterations with early exaggeration: 68.492363
[t-SNE] Error after 2000 iterations: 0.961622
```

Now we have two-dimensional representations ready !! lets plot the clusters using Bokeh. It will be useful to derive the centroid location of each topic before that. This would add context to the visualization.

```
In [58]: # Define helper functions
def get_mean_topic_vectors(keys, two_dim_vectors):
    '''returns a list of centroid vectors from each predicted topic category'''
    mean_topic_vectors = []
    for t in range(n_topics):
        articles_in_that_topic = []
        for i in range(len(keys)):
            if keys[i] == t:
                articles_in_that_topic.append(two_dim_vectors[i])

        articles_in_that_topic = np.vstack(articles_in_that_topic)
        mean_article_in_that_topic = np.mean(articles_in_that_topic, axis=0)
        mean_topic_vectors.append(mean_article_in_that_topic)
    return mean_topic_vectors
```



```
In [59]: from bokeh.plotting import figure, output_file, show
from bokeh.models import Label
from bokeh.io import output_notebook
output_notebook()

colormap = np.array([
    "#1f77b4", "#aec7e8", "#ff7f0e", "#ffbb78", "#2ca02c",
    "#98df8a", "#d62728", "#ff9896", "#9467bd", "#c5b0d5",
    "#8c564b", "#c49c94", "#e377c2", "#f7b6d2", "#7f7f7f",
    "#c7c7c7", "#bcbd22", "#dbdb8d", "#17becf", "#9edae5" ])
colormap = colormap[:n_topics]
```

(<https://bokeh.pydata.org>) successfully loaded.


Now lets plot the cluster and see..Also we have included top 3 words in the cluster, that are located on the centroid of that topic.

```
In [61]: top_3_words_lsa = get_top_n_words(3, lsa_keys, small_document_term_matrix, small_count_vectorizer)
lsa_mean_topic_vectors = get_mean_topic_vectors(lsa_keys, tsne_lsa_vectors)

plot = figure(title="t-SNE Clustering of {} LSA Topics".format(n_topics), plot_width=700, plot_height=700)
plot.scatter(x=tsne_lsa_vectors[:,0], y=tsne_lsa_vectors[:,1], color=colormap[lsa_keys])

for t in range(n_topics):
    label = Label(x=lsa_mean_topic_vectors[t][0], y=lsa_mean_topic_vectors[t][1],
                  text=top_3_words_lsa[t], text_color=colormap[t])
    plot.add_layout(label)

show(plot)
```



As we can see above plot is not very convincing. not much separation of topics is achieved , and it is difficult to tell whether this can be attributed to the LSA decomposition or instead the t -SNE dimensionality reduction process.

Latent Dirichlet Allocation

Below process is same. Just that, instead of using LSA, I used LDA. LDA is instead a generative probabilistic process, designed with the specific goal of uncovering latent topic structure in text corpora.

```
In [62]: from sklearn.decomposition import LatentDirichletAllocation

lda_model = LatentDirichletAllocation(n_components=n_topics, learning_method=
'online',
                                     random_state=0, verbose=0)
lda_topic_matrix = lda_model.fit_transform(small_document_term_matrix)
```

```
In [63]: lda_keys = get_keys(lda_topic_matrix)
lda_categories, lda_counts = keys_to_counts(lda_keys)
```

```
In [64]: top_n_words_lda = get_top_n_words(10, lda_keys, small_document_term_matrix, sm
all_count_vectorizer)

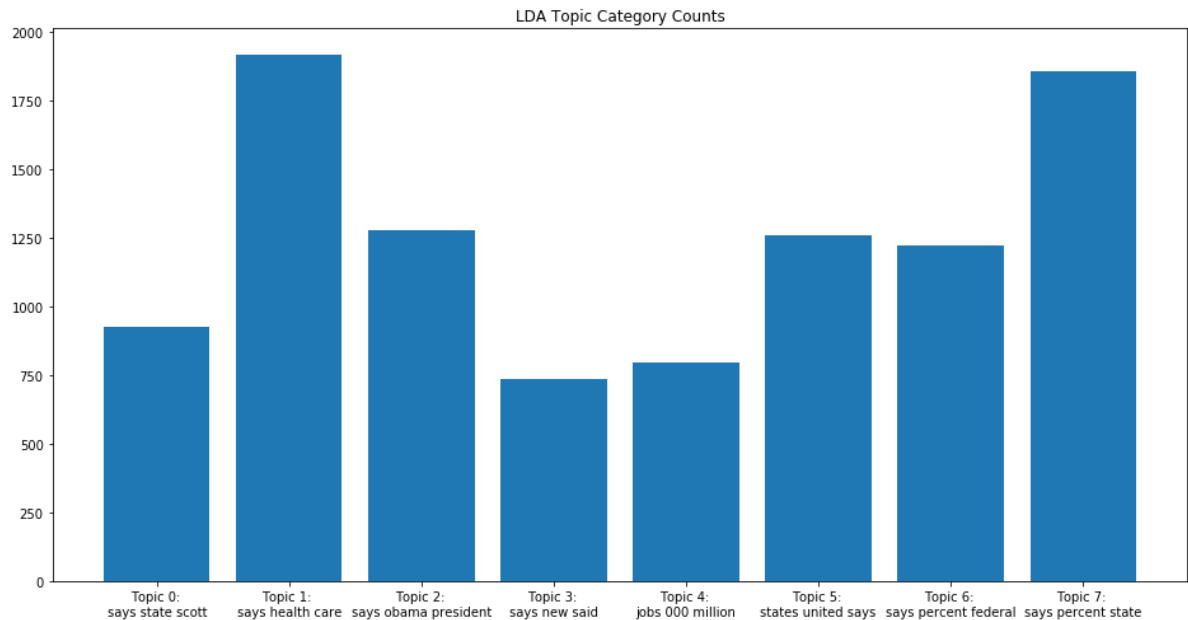
for i in range(len(top_n_words_lda)):
    print("Topic {}: ".format(i), top_n_words_lda[i])
```

```
Topic 0:  says state scott gov federal rick government court wisconsin walker
Topic 1:  says health care percent tax taxes people voted obama plan
Topic 2:  says obama president barack clinton said hillary trump donald campa
ign
Topic 3:  says new said obama state administration water make government texa
s
Topic 4:  jobs 000 million new says job health created insurance sector
Topic 5:  states united says rate percent years year country state people
Topic 6:  says percent federal people year state illegal tax american america
ns
Topic 7:  says percent state 000 billion tax year years budget texas
```

```
In [65]: top_3_words = get_top_n_words(3, lda_keys, small_document_term_matrix, small_count_vectorizer)
labels = ['Topic {}: \n'.format(i) + top_3_words[i] for i in lda_categories]

fig, ax = plt.subplots(figsize=(16,8))
ax.bar(lda_categories, lda_counts)
ax.set_xticks(lda_categories)
ax.set_xticklabels(labels)
ax.set_title('LDA Topic Category Counts')
```

Out[65]: Text(0.5,1,'LDA Topic Category Counts')



To achieve better comparison, we again take this topic matrix and project it into two dimensions with t-SNE.

```
In [66]: from sklearn.manifold import TSNE

tsne_lda_model = TSNE(n_components=2, perplexity=50, learning_rate=100,
                      n_iter=2000, verbose=1, random_state=0, angle=0.75)
tsne_lda_vectors = tsne_lda_model.fit_transform(lda_topic_matrix)


[t-SNE] Computing 151 nearest neighbors...
[t-SNE] Indexed 10000 samples in 0.016s...
[t-SNE] Computed neighbors for 10000 samples in 1.522s...
[t-SNE] Computed conditional probabilities for sample 1000 / 10000
[t-SNE] Computed conditional probabilities for sample 2000 / 10000
[t-SNE] Computed conditional probabilities for sample 3000 / 10000
[t-SNE] Computed conditional probabilities for sample 4000 / 10000
[t-SNE] Computed conditional probabilities for sample 5000 / 10000
[t-SNE] Computed conditional probabilities for sample 6000 / 10000
[t-SNE] Computed conditional probabilities for sample 7000 / 10000
[t-SNE] Computed conditional probabilities for sample 8000 / 10000
[t-SNE] Computed conditional probabilities for sample 9000 / 10000
[t-SNE] Computed conditional probabilities for sample 10000 / 10000
[t-SNE] Mean sigma: 0.090871
[t-SNE] KL divergence after 250 iterations with early exaggeration: 85.454643
[t-SNE] Error after 2000 iterations: 1.739634
```

```
In [67]: top_3_words_lda = get_top_n_words(3, lda_keys, small_document_term_matrix, small_count_vectorizer)
lda_mean_topic_vectors = get_mean_topic_vectors(lda_keys, tsne_lda_vectors)

plot = figure(title="t-SNE Clustering of {} LDA Topics".format(n_topics), plot_width=700, plot_height=700)
plot.scatter(x=tsne_lda_vectors[:,0], y=tsne_lda_vectors[:,1], color=colormap[lda_keys])

for t in range(n_topics):
    label = Label(x=lda_mean_topic_vectors[t][0], y=lda_mean_topic_vectors[t][1],
                  text=top_3_words_lda[t], text_color=colormap[t])
    plot.add_layout(label)

show(plot)
```



Now that impressive! Its evident that LDA has had much more succcess than LSA in distinguishing the topic categories. Therefore, LDA is an appropriate algorithm for our case.